

# Algorithms that Remember: Model Inversion Attacks and Data Protection Law

Michael Veale\*, Reuben Binns† and Lilian Edwards‡

\* University College London, United Kingdom [m.veale@ucl.ac.uk]

† University of Oxford, United Kingdom [reuben.binns@cs.ox.ac.uk]

‡ University of Strathclyde, United Kingdom [lilian.edwards@strath.ac.uk]

*Draft version May 22, 2018*

## Abstract

Many individuals are concerned about the governance of machine learning systems and the prevention of algorithmic harms. The EU’s recent General Data Protection Regulation (GDPR) has been seen as a core tool for achieving better governance of this area. While the GDPR does apply to the use of models in some limited situations, most of its provisions relate to the governance of personal data, while models have traditionally been seen as intellectual property. We present recent work from the information security literature around ‘model inversion’ and ‘membership inference’ attacks, which indicate that the process of turning training data into machine learned systems is not one-way, and demonstrate how this could lead some models to be legally classified as personal data. Taking this as a probing experiment, we explore the different rights and obligations this would trigger and their utility, and posit future directions for algorithmic governance and regulation.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>European data protection law and machine learning</b>	<b>3</b>
2.1	Limits of the harms addressed by the current regime . . . . .	3
2.2	Models on the move . . . . .	4
<b>3</b>	<b>Why might models be personal data</b>	<b>6</b>
3.1	Attack types . . . . .	6
3.1.1	Model inversion . . . . .	7
3.1.2	Membership inference . . . . .	8

3.2	Why these attacks render models as personal data . . . . .	8
<b>4</b>	<b>Implications of models as personal data</b>	<b>10</b>
4.1	For data subjects . . . . .	11
4.1.1	Information rights . . . . .	11
4.1.2	Erasure rights . . . . .	12
4.2	Right to restrict processing and to object . . . . .	13
4.3	For data controllers . . . . .	14
4.4	Security principle and data protection by design . . . . .	14
4.5	Storage limitation principle . . . . .	16
<b>5</b>	<b>Discussion and conclusion</b>	<b>16</b>

## 1 Introduction

A recent heated topic in technology law and policy is whether machine learning systems are sufficiently regulated, a worry compounded by their apparent potential to reproduce societal discrimination and to transform mundane personal data into sensitive, invasive insight. The recent EU General Data Protection Regulation (GDPR), which strengthens data protection provisions and penalties, has been looked to internationally as a way forward, particularly following high-profile coverage of the scandal around Facebook and Cambridge Analytica in early 2018. Yet given that data protection law applies primarily to personal data, not aggregated analytic products resulting from it, the extent to which is is normally thought to provide control over models themselves is limited [27]. Instead, these have been typically thought to be primarily governed by varying intellectual property rights such as trade secrets, and typically discussed as such [16].

In this paper, we argue that changing technologies may render the situation may be less clearcut. We highlight a range of recent research which indicates that the process of turning training data into machine learned systems is not one-way but two: that training data, a semblance or subset of it, or information about who was in the training set, can in certain cases be reconstructed from a model. The consequences of this for the regulation of these systems is potentially hugely significant, as the rights and obligations for personal data differ strongly from the few generally thought applicable to models.

First, we introduce relevant element of data protection law in this context for a broad audience in the context of current debates around algorithmic systems and of the growing trends to trade and enable access to models, rather than underlying datasets. Second, we introduce model inversion and membership inference attacks, describing their setup and why data protection law would be likely to classify vulnerable models as personal data. We then describe selected consequences of this change for data subjects, who would have access to new information, erasure and objection rights as a result, and for the modellers and model

recipients, who would need to consider security, ‘data protection by design’ and storage limitation implications. We conclude with a short discussion about the utility of this approach, and reflections on what this setup tells us about desirable directions for future lawmaking around fast-moving technological practices.

## 2 European data protection law and machine learning

The GDPR is a lengthy and complex law, and is challenging to concisely summarise. It applies whenever personal data is processed (including collected, transformed, consulted or erased) either within the Union, or outside of the Union relating to an EU resident. Personal data is defined by how much it can render somebody identifiable—going beyond email or phone number to include dynamic IP addresses, browser fingerprints or smart meter readings. The individual data relates to is called the *data subject*. The entities who determine ‘the purposes and means of the processing of personal data’ are *data controllers*. Data subjects have rights over personal data, such as rights of access, erasure, objection to processing, and portability of their data in a common machine-readable format elsewhere. Data controllers are subject to a range of obligations, such as ensuring confidentiality, notifying if data is breached, and undertaking risk assessments. Additionally, they must only process data where they have a lawful basis—such as consent—to do so, for a specified and limited purpose, and a limited period of storage.

Data protection law already governs collection and use of data in generating machine learning models, and in certain limited conditions, the application of model results upon data subjects. Models cannot be trained from personal data without a specific lawful ground, such as consent, contract or legitimate interest; data subjects should be informed of the intention to train a model, and usually maintain a right to object or withdraw consent; and in situations where models inform a significant, solely automated decision, individuals can appeal to the data controller for meaningful information about the logic of processing, or to have the decision taken manually reconsidered. The use of machine learning to turn ‘normal’ personal data into ‘special category’ personal data, such as race, political opinion, or data concerning health, also requires establishing a lawful basis, which will usually be more stringent than personal data in general. The nature and utility of these provisions is explored and debated extensively elsewhere [51, 20, 42, 28]; for reasons of space we do not expand upon them here.

### 2.1 Limits of the harms addressed by the current regime

If data protection already implicates machine learning in the above ways, does it matter if a model is personal data or not? In short—yes. The GDPR does touch upon machine learning, but largely indirectly, governing systems only when personal data is involved in training them or querying them to apply their results.

Yet provisions would have a significantly different effect were models themselves to benefit from some of the status that personal data has. As it stands, there are no data protection rights nor obligations concerning models in the period after they have been built, but before any decisions have been taken using them. Even the provisions that apply outside this period are relatively minor: those models being used as decision-support, rather than as decision-making instruments, or those significantly affecting groups rather than individuals, are subject to few solid protections at all [20].

There is reason to believe this level of control is insufficient. Individuals might want to control how they specifically are ‘read’ by machine learned systems [28]. Trained models can transform seemingly non-sensitive data, such as gait or social media use, into sensitive data, such as information on an individual’s fitness or medical conditions. In many fields, these models are far from user-independent, and so individuals adding their own granular training data to these models enables them to predict more accurately in the future. Take automated lipreading systems. These transform video of speech into approximated transcripts, and, being largely speaker-dependent [7], require individually specific training data to be effective. Once a model integrates an individual’s data, it will predict their speech with significantly greater accuracy than that of others. Do individuals have a right to reverse this process, and have some agency over the model after it has been trained and potentially traded? Do they even have some right to utilise the model, or at least to know what it is used for, given their stake in training it?

In a similar vein, arguments have been forwarded that ‘groups’ of individuals deserve agency over their representation in models. Scholars concerned with ‘categorical’ or ‘group privacy’ see such groups as having collective rights to determine how their identities are constituted [50, 46]. Within this view, an ethical breach might be considered to have occurred “when data or information is added to subject’s identity without consent” [35]. For example, given claims of correlations between smartphone-captured data and aspects of physical/mental health [21], individuals revealing a rare condition to a modeller might enable such inferences to be successfully made for others, too. Similar arguments also exist in relation to data that connects individuals more deeply, such as genomic data, where sequencing the genome of one family member might reveal information about many.<sup>1</sup>

## 2.2 Models on the move

These issues are of increasing importance given how data controllers increasingly refrain from trading data, as the ability to do this freely is heavily limited by

---

<sup>1</sup>A related concern surrounds what one might call ‘fair trade algorithms’—that using a system with data that has been collected unethically, perhaps in a country with limited fundamental rights or privacy rules, is itself unethical. As the users of such a system are unlikely to be in the training set, we do not address this issue here.

data protection law, and instead are looking to trade or rent out models trained on it, as a way to pass on the value with fewer privacy and regulatory concerns. Many large firms already offer trained models for tasks including face recognition, emotion classification, nudity detection and offensive text identification. Two main business models underpin this.

The first is in the **licensing of APIs** through “App Store”-like platforms. Firms earn royalties when their models are deployed. Microsoft’s *Cortana Intelligence Gallery* and Algorithmia’s *Marketplace* match API providers with potential customers. Google’s *Cloud AutoML* uses transfer learning to allow firms to enrich Google’s generic models with their own specialised datasets. The trend towards augmentable, pre-trained “learnware” [53] appears to be accelerating.

The second is the trading of **packaged models**. This might be preferable where APIs over the Internet are too sluggish, where queries are highly sensitive (e.g. medical records), or where transparency requirements require full access to model specifications. Apple’s recent *Core ML* library and Google’s *Tensorflow Mobile* are designed to run pre-trained models, some of which they provide, on portable devices.

In relation to many aspects of privacy, model trading is an attractive prospect. It stands in stark contrast to the incumbent (and, many suspect, largely illegal) system of large scale, indefinite data accumulation and retention by shadowy and distant data brokers—systems presenting significant difficulties for data subject comprehension and control [48, 8]. While it is commonly claimed that the “free flow” of data is of economic benefit, much of this benefit derives from the movement of mined insights rather than the transmission of individual records. Given that the push towards open data might, in at least some regards, be at tension with privacy [5], model trading might serve as a useful approach to balance trade-offs that emerge.

Furthermore, model trading might mitigate concerns around platform monopolies. Enabling users to deploy local personalisation tools might balance power relations in relation to large firms hoarding personal data. This is the vision of *personal data container* proponents, who see distributed local processing, powered by decentralised, privacy-preserving analytical techniques—such as secure multi-party computation and homomorphic encryption—as enabling a shift of economic incentives and control from large players back to data subjects [14]. Many processes that use highly granular knowledge for profiling, such as the delivery of advertisements or the personalisation of news media, are easily envisaged as happening using edge computing, rather than in the cloud, limited the sensitive data leaving devices a user directly controls [15].

Yet such systems, no longer representing unique records which might render an individual identifiable, have not been considered as personal data, and have thus conventionally been excluded from the data protection regime. In this paper, we challenge that conventional understanding, and reflect upon the legal provisions this would trigger. Recent evidence, reviewed here, highlights that models themselves may leak data they were trained with—raising classic data

confidentiality concerns. Data protection rights and obligations might then apply to models themselves. We next outline the format model inversion attacks can take, why they might render models as personal data in the sense of European protection law, and what the consequences for this might be for data subjects and for data controllers.

### 3 Why might models be personal data

It has been demonstrated that machine learning models are vulnerable to a range of cybersecurity attacks that cause breaches of confidentiality. Confidentiality attacks leak information to those other than whom designers intended to view it. In the case of machine learning systems, there are different types of these attacks. The first concerns *model stealing*, e.g. where an attacker uses API access to replicate a model [47]. Without a further confidentiality breach, this is primarily a concern for intellectual property rather than privacy, and of less concern here. A second attack class, *model inversion*, turns the journey from training data into a machine learned model from a one-way one to a two-way one, permitting the training data to be estimated with varying degrees of accuracy. A third attack class, *membership inference*, does not recover the training data, but instead recovers information about whether a particular individual was in the training set or whether they were not. Both model inversion and membership inference can be undertaken as a *black box attack*, where the attack can be done with only query access (e.g. through the API business model above), or an *white box attack*, where an attacker requires full access to the model’s structure and parameters [22, 1].

We will formally describe both model inversion and membership inference attacks in a manner amenable to the discussion of personal data. The setup is as follows. A data controller holds a dataset of personal data  $\mathbf{A} = a_{i,j} \in \mathbb{R}^{m \times n}$  where each  $[a_{1,*}, a_{2,*}, \dots, a_{m,*}]$  within is a row of personal characteristics relating to one of the  $m$  data subjects in the set  $|DS_1| = m$ , with each of the  $n$  variables indexed by  $j$ . They also have access to a model  $M(\mathbf{B})$ , which is a machine learned model trained on personal data  $\mathbf{B} = b_{i,j} \in \mathbb{R}^{x \times y}$  where each  $[b_{1,*}, b_{2,*}, \dots, b_{x,*}]$  within is a row of personal characteristics relating to one of  $x$  data subjects in the set  $|DS_2| = x$ , and each one of the  $y$  variables a feature in the trained model. The data controller may have access to the model either directly (white-box) or via a query interface (black-box). We assume  $DS_1 \cap DS_2 > 0$ : that is, some individuals are in both the training set and the additional dataset held. We refer to individuals in both  $\mathbf{A}$  and  $\mathbf{B}$  as in set  $Z$ .

#### 3.1 Attack types

In the following section, we outline, at a high level and non-exhaustively, forms and work on model inversion and membership inference attacks. A visual depiction of both can be found in Figure 1.

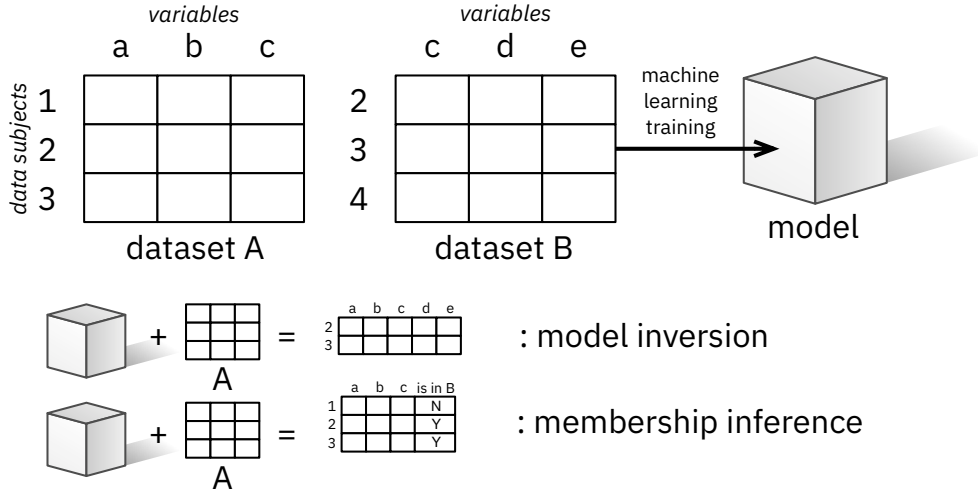


Figure 1: Model inversion and membership inference attacks. Diagram by authors.

### 3.1.1 Model inversion

Under a model inversion attack, a data controller who does not initially have direct access to **B**, but is given access to **A** and  $M(\mathbf{B})$ , is able to recover some of the variables in training set **B**, for those individuals in both the training set and the extra dataset **A**. These variables connect to each other, such that the new personal dataset in question has all the variables of **A** and some of **B**. There may be error and inexactitude in the latter, but the data recovered from those in the training dataset will be more accurate than characteristics simply inferred from those that were not in the training dataset.

One of the earliest attacks resembling model inversion emerged from the field of recommender systems; a demonstration that collaborative filtering systems, where item recommendations are generated for a user based on behavioural patterns of other users, can end up revealing the consumption patterns of individual users [10]. At the time, a reader might have reasonably assumed such risks to be a quirk of the particular system and application area, but subsequent work suggests that this may be a more general problem potentially facing any kind of ML model that utilises personal data for training. Later work undertook attacks against several machine learning models to attempt to learning meaningful information about the training data. They were not concerned with privacy *per se* but with trade secrets, seeking to uncover the “secret sauce” of algorithms that might give them a commercial edge, such as whether speech recognition systems were trained on certain accents or not [3]. [23] examine models designed to select correct medical doses for a widely used anticoagulant that interacts

strongly with individual genetic markers. They show the possibility of reverse-engineering to reveal patients’ genetic markers with some demographic information about patients in the training data.<sup>2</sup> Further work demonstrated both white- and black-box attacks re-identifying individuals from models trained on survey data with no false positives, and black-box attacks to reconstruct faces from facial recognition systems to the point where skilled crowdworkers could use the photo to identify an individual from a lineup with 95% accuracy [22].

Other work connected to model inversion has indicated that only small changes to training algorithms lead to nearly indistinguishable models that are possible to exploit to leak large amounts of private data [45], or that systems exactly memorise specific private information in training sets, such as strings of sensitive data [12]. Some model structures also require the training data in order to function, such as support vector machines, where the vectors that define the classification boundary, drawn from the training data, are bundled with the model used.<sup>3</sup>

### 3.1.2 Membership inference

Membership inference attacks do not recover training data, but instead ascertain whether a given individuals’ data was in a training set or not. Under a membership inference attack, the holder of  $\mathbf{A}$  and  $M(\mathbf{A})$  does not recover any of the columns in  $\mathbf{B}$ , but can add an additional column to dataset  $\mathbf{A}$  representing whether or not a member of  $DS_1$  is in the set  $Z$ : that is, whether or not they were also part of the training set participants  $DS_2$ .

[43] demonstrate membership inference in a black box attack against a hospital discharge model. Connectedly, [41] looked at membership inferences in location data, showing that it is comparatively easy to examine whether an individual with a certain movement pattern was used in the construction of an aggregate. They note this could be concerning if such aggregates were themselves made to understand a sensitive subgroup, such as individuals with dementia.

## 3.2 Why these attacks render models as personal data

While these attacks are nascent, their potential is being increasingly demonstrated.<sup>4</sup> The initial important question for data protection law is to what extent would systems vulnerable to attacks like these be considered datasets of personal data?

There are strong arguments that a model inversion–vulnerable  $M(\mathbf{B})$  would be personal data. A direct analogy can be made to personal data which has been

<sup>2</sup>Some argue that this paper shows no privacy breach here beyond people finding machine learning “creepy”: see [33].

<sup>3</sup>Thanks to REMOVED FOR BLIND REVIEW for bringing this point to our attention in this context.

<sup>4</sup>We are, however, cautious about overstating their practical efficacy, which does remain unclear. We are not aware of any documented attacks “in-the-wild”.



‘pseudonymised’. Article 4(5) of the GDPR defines pseudonymisation as ‘the processing of personal data in such a manner that the personal data can no longer be attributed to a specific data subject without the use of additional information, provided that such additional information is kept separately and is subject to technical and organisational measures to ensure that the personal data are not attributed to an identified or identifiable natural person’. Under European law, pseudonymised data explicitly remains personal data. In the above setup,  $M(\mathbf{B})$  represents the pseudonymised version of the training set  $\mathbf{B}$ , while  $\mathbf{A}$  represents the key by which such data can be partially reidentified. Where a single data controller is in possession of both  $\mathbf{A}$  and  $M(\mathbf{B})$ , and a model inversion attack is possible, it would appear by analogy that not only  $\mathbf{A}$  but also  $M(\mathbf{B})$  should be considered personal data.

Of potentially greater interest however is the situation where a model  $M(\mathbf{B})$  has been released, and so  $\mathbf{A}$  and  $M(\mathbf{B})$  are held by different entities. There is legal precedent for the model  $M(\mathbf{B})$  to be considered personal data in this case too.

Personal data means any information relating to an identified or identifiable natural person (art 4(1)). The surprising ease with which individuals can be identified in supposedly anonymous datasets [38] creates a great deal of uncertainty around which datasets are not personal data [40]. Recital 26 provides an aide to navigating this problem, with a test of reasonable likelihood of reidentification.

To determine whether a natural person is identifiable, account should be taken of all the means reasonably likely to be used, such as singling out, either by the controller or by another person to identify the natural person directly or indirectly. To ascertain whether means are reasonably likely to be used to identify the natural person, account should be taken of all objective factors, such as the costs of and the amount of time required for identification, taking into consideration the available technology at the time of the processing and technological developments.

When asked in recent years, the Court of Justice of the European Union (CJEU) has set a low bar and wide scope for what is considered personal data. In *Breyer* (ECLI:EU:C:2016:779), the CJEU clarified the reasonable likelihood test in the Data Protection Directive 1995 that now forms part of Recital 26 of the GDPR. In particular, the Court took a wide view of identifiability, clarifying that where additional data is required to reidentify a dataset, that data need not all be in the hands of a single entity to be considered personal data. In effect, this illustrates that in the case of either attack  $M(\mathbf{B})$  might be considered personal data if dataset  $\mathbf{A}$  is held by another entity.

Furthermore, while the personal data in model inversion attacks is quite easily construed as personal data, insofar as it resembles a training set or can be used to identify individuals, it might initially appear less clear that data obtained from membership inference—whether an individual was in a training dataset—is

personal data. Personal data does not specify particular sensitive or private categories in the definition, instead using the phrase ‘any information’. In the recent CJEU case *Nowak* (ECLI:EU:C:2017:994), the Court engaged with the meaning of this phrase, noting that it is “not restricted to information that is sensitive or private, but potentially encompasses all kinds of information, not only objective but also subjective”. Similarly, and following the Article 29 Working Party, the pan-European group of regulators that worked together to provide guidance on the Data Protection Directive, the court understood that this “any information” can “relate to” an individual in many ways: by content, by purpose and by effect. Seen through this lens, information on an individuals’ membership of a training set would indeed fall within the scope of personal data, regardless of how trivial or mundane it might be to the individual it concerns.

There does exist an argument forwarded by scholars such as Purtova [40] which argues that the Court’s approach to the scope of personal data fuels undesirable data protection maximalism. She provocatively claims where common environmental factors which do not identify an individual by themselves, such as the weather, are used in a smart city context as an input to a behavioural profiling system. In this case, she argues, it is plausible that the weather data would be personal data by means of purpose and effect. In the context of this paper, her argument might plausibly be extended to claim that the weights, and perhaps the structure, of machine learning models relate to an individual by means of impact, and by virtue of this are personal data. While we acknowledge this *reductio ad absurdum* argument concerning the current scope of personal data, and the consequences for it as making the law impracticably broad, our argument does not lean in this direction. We do not aim to support, oppose or resolve the dilemma raised by Purtova; but merely to note that the argument made here—that inverted models might fall under the definition of personal data—does not depend on the kind of expansive definition that gives rise to such absurdities. Thus, even if the definition were to be somehow tightened in scope (indeed, the scope of personal data has changed even between recent cases such as *YS and others* (ECLI:EU:C:2014:2081) and *Breyer*), the argument above concerning inverted models would still likely stand.

In sum, model inversion and membership inference attacks, where possible, do risk models being considered as personal data even without resorting to a maximalist reading of data protection law. The question then remains—what are the practical consequences of this, and are they of interest to those hoping to better control these systems and hold them to account?

## 4 Implications of models as personal data

Were models to be personal data, very different sets of rights and obligations would trigger compared to their status today. In this section, we survey a salient range of them from two angles—rights that would be newly afforded to the data

subject, and obligations placed on the data controller.

## **4.1 For data subjects**

We consider three tasks data subjects might want to achieve in relation to models trained with their data: where a data subject wishes to access models and to know where they have originated from and to whom they are being traded or transmitted; where a data subject wishes to erase herself from a trained model; and where a data subject wishes such a model not to be used in the future.

### **4.1.1 Information rights**

The GDPR contains information provisions which are triggered on request, such as the right of access (art 15) and the right of portability (art 20), as well as provisions where the data controller must provide information to the data subject without being solicited to do so (arts 13–14). In practice, these rights are most well-known for the requirement, carried through from the Data Protection Directive, to provide all the information held about an individual upon request. It seems unlikely that such a request would allow an individual to request an entire model on the basis that it comprised in part of their data as, except in the rare case where only a single individuals’ data was used for the model building,<sup>5</sup> to do so might comprise of the privacy of others in the training set. To provide it in that way would be a breach of data protection’s security principle. It also seems unlikely that they would be successful with much ease at requesting a copy of the data retrieved from a model, as it is unlikely to include a name or identifier, and to make names and identifiers difficult to access is a common practice by data controller which has the effect of rendering personal data difficult to exercise rights over [49].

Particularly given the increasing interest in model trading described above, it is more interesting to consider other aspects of these information rights. In particular these rights might enable better tracking of the provenance of data, analysis and decision-making systems, something which has received increasing attention of late [44]. In effect, these rights help track the origin and destination of trained models. The GDPR does have two core types of requirement which respectively require controllers with data to provide information about from whom specifically personal data came from, and to whom, more generally, it is going.

At the point of collection (art 13(1)(e)), and upon request (art 15(1)(c)), the data subject should be provided with ‘the recipients or categories of recipients of the personal data, if any’. While data protection authorities interpretations have indicated that providing potentially broad and unhelpful ‘categories’ alone is in line

---

<sup>5</sup>Such a case might include, for example, a tailored voice recognition system or a learned smart home personalisation model. Portability of these models to another controller might be useful, although it would introduce a likely infeasible compatibility demand on the receiving provider, who would likely find it easier simply to take the raw data and retrain a model themselves.

with the law [24], recent guidance from the Article 29 Working Party notes that [2] only providing the ‘categories’ will require a justification of why this is in line with the principle of fairness in data protection, and if it is done, such categories must ‘be as specific as possible by indicating the type of recipient (i.e. by reference to the activities it carries out), the industry, sector and sub-sector and the location of the recipients.’

Where one controller receives personal data from another—as is this case with model trading if we conceive of models that way, there is an alternative approach afforded by Article 14(2)(f). This provision states that ‘the controller shall provide the data subject with the following information necessary to ensure fair and transparent processing in respect of the data subject [...] from which source the personal data originate, and if applicable, whether it came from publicly accessible sources’.

The only applicable exemption here is Article 14(5)(b) where ‘the provision of such information proves impossible or would involve a disproportionate effort’. It is highly likely that a controller faced with the situation of model-as-personal-data would try to utilise this exemption, particularly as they are unlikely to have contact information for the data subjects in the model, and Article 11(1) notes they are not obliged to hold such additional data for the sole purposes of complying with the GDPR. This is however not the end of the story, because Article 14(5)(b) goes on to note that where disproportionate effort applies, ‘[in] such cases the controller shall take appropriate measures to protect the data subject’s rights and freedoms and legitimate interests, including making the information publicly available.’ The one safeguard that is explicitly listed would seem to imply that the data controller might then be obliged to make publicly available—perhaps through their website—their receipt of a model and the sources from which the personal data originate.

In effect, this would mean that organisations receiving invertible or potentially invertible models would have to publish where they came from. To do so would be a huge asset to those mapping these flows and looking to understand the data economy and algorithmic accountability within it, particularly if such notices were machine readable in form.

#### **4.1.2 Erasure rights**

The famous ‘right to be forgotten’ is a qualified right of a data subject to “the erasure of personal data concerning him or her” (Article 17). Core reasons a data subject might want to erase herself from a model overlap with the general reasons for model control presented earlier—to erase insights about her she might dislike; to erase unwanted insights about a group she identifies as part of; or to erase insights which might lead to data breaches.

There are two main ways to erase data from a trained model. Firstly, a model can be trained based upon an amended training dataset. The computational intensity of much machine learning training, even on the world’s most powerful

computational infrastructures, means such training is far from cost-free or instantaneous. Retraining brings significant energy costs—data centres consume between 1.1–1.5% of global power consumption [32], as well as time and labour costs.

Secondly, the model itself can be amended after training. This is not easy, and rarely currently possible in modern systems. Approaches for quick and easy “machine unlearning” are only beginning to be proposed and are still largely unexplored, let alone at a stage ready for deployment [11, 6]. Currently proposed methods cannot be retrofitted onto existing systems, and would require entire model pipelines to be reconceived, with unclear effects. Consequently, large-scale organisations such as Google do not appear to remove links from their trained search model, but filter (“delist”) results between model output and search result delivery [26].

Yet while it might be possible for a single individual to remove a pattern from a system trained on a great deal of data from a few individuals, where the number of observations in a dataset is similar to the number of data subjects, it would appear difficult for an individual to remove such patterns alone (except from poorly trained models). Models basing patterns on single data records are generally considered to have been *overfitted*—memorising the training data rather than finding generalisable patterns.<sup>6</sup> Such a right becomes more useful when used collectively. Whether this is possible will depend both on the phenomena at hand, and the level of coordination between the data subjects. Co-ordination in the form of online petitions and crowdfunding make this an interesting avenue to explore for the right of erasure, and one that is already being probed in relation to the right of access [31].

## 4.2 Right to restrict processing and to object

A set of important rights that have received little media attention essentially surround the rights of individuals to say no to types of processing which they did not consent for. If models were personal data, the consequences of this are quite difficult to say. The right to object in Article 21 allows an individual to object to processing based on the legitimate interests of the data controller of public interest (largely public sector processing), as long as the data controller cannot demonstrate a ‘compelling’ legitimate interest of their own—something which seems a high bar indeed, and seemingly unlikely to be met unless some social benefit is present, rather than just controllers’ economic incentive.

The right to restrict processing in Article 18 has a wider range of provisions, broadly giving the data subject rights to stop processing of their data until, for

---

<sup>6</sup>There are times when one record can make a difference: “one-shot learning” in image recognition, for example. As children, humans are very good at seeing an object once, classifying it, and then correctly identifying a second, different object of the same classification is presented. Standard machine learning systems perform very poorly at this task, requiring many examples of a single class in order to become good at recognising it.

example, a balancing test relating to the right to object can be carried out, or until the accuracy of the data is ascertained and rectified. These are time limited, and while there is no balancing test to use them (and therefore they could be used quite disruptively), they are generally considered lesser.

What is it to object to the use of a model? To query a model might be considered analogous to consulting the personal dataset within. Consultation is one of the many actions that explicitly comprise processing in data protection. Yet because the data is not organised by record when it is in the form of the model, querying a model seems more like querying the entire dataset than querying a single record. Does objection give every individual a veto right over the consultation of the model in its entirety, for any lawful purpose? This would seem highly problematic, and likely disproportionate, but it is a possible reading of this provision, and one which demonstrates the difficulties seeing models as personal data would engender.

A more sensible approach from a user perspective to achieve their goals would often be to restrict the processing of that model in relation to an individual decision, but would be already possible by using these rights to object or restrict the personal data constituting the query being used for prediction in a specific case, or possibly the right not to be subject to solely automated decision-making (Article 22) which, unusually for data protection, instead target the decisions or measures themselves rather than the use of data. Yet the strange interplay between objection and models made of mixed-up personal data could potentially be a place of tension were individuals to try to test and enforce these rights.

### **4.3 For data controllers**

Data controllers must consider both the specific rights and obligations they are subject to, as well as adherence to the overarching principles of data protection. We highlight two relevant areas here in the context of models as personal data: the security principle and the storage limitation principle.

### **4.4 Security principle and data protection by design**

As an overarching concern and obligation, data controllers need to consider whether their system leaks personal data in the first place. It seems unlikely that a modeller would wish to establish a legal basis for onward transfer of a model they have trained, something which would be especially onerous if that model is being transmitted outside of the EU or a country deemed ‘adequate’ with EU data protection law. If they do transfer such a model without a legal basis to do so, and such a model is inverted, it would likely be considered both a data breach and a violation of the security principle more generally. This is a relatively shocking conclusion, and endangers many of the positive aspects of model trading, particularly that large datasets are not being transmitted constantly, which this paradigm promises.

Further reason to ensure that models are not invertible comes from Article 25(1) of the GDPR, which introduces a qualified obligation to implement technical and organisational measures designed to implement data protection principles. In combination, these aspects require us to consider how models can be made which are resilient to the attacks described above.

Thankfully, security researchers try to secure systems as well as try to break them. The most common defence that can be levied against model inversion attacks is differential privacy [18]: a defence discussed in most of the papers above describing the threats. Differential privacy is a formalised notion of privacy as information disclosure.<sup>7</sup> At the core of differential privacy is the notion of adding precise forms of random noise to queries such that that for every individual in a dataset, removing them from that dataset will not noticeably change the results of that query.<sup>8</sup> Differential privacy guarantees essentially limit what can be learned about any individual in the dataset to that which could be learned about them from everybody else’s data alone. Differentially private approaches to machine learning work in a similar way: models lacking data relating to a particular individual function extremely similarly to those containing it [29]. This provides intrinsic protection against a great deal of model inversion attacks.

*Theoretically*, such learning algorithms are just as powerful as non-differentially private ones [30]. Since algorithms are supposed to find generalisable patterns, not memorise and regurgitate specific records, this makes intuitive sense (discussed further below). Yet despite the growing interest in differential privacy [39, 17, 9], the real challenge comes with deployment. The tools available today can be computationally expensive to deploy [13] as well easily undermined with even small or arcane software errors [34]. Only a few large and powerful companies have demonstrated an ability to deploy them, and only then for very limited purposes. Furthermore, differential privacy works well at protecting disclosure in contexts where every individual has the same weight, such as in a count, but poorly in situations where there are extreme outliers. In a dataset of wealth containing Bill Gates, as the amount of noise needed to be added reduces the query results to absurdity [36, 4]. In cases where outliers might be vulnerable as well as powerful this is problematic: we might suspect it is they who need data protection the most [20].

A second linked line of defence, albeit without the guarantees differential privacy provides, is to attempt to make models that do not “overfit” to the data. Where they do, they are confusing the “noise” in the dataset for the generalisable “signal” that helps prediction on unseen cases: memorising rather than doing anything close to learning. Avoiding overfitting is important—and in this way, data protection by design might legally oblige the training of methodologically sound models—but avoiding it is not enough to guarantee a model will not be

---

<sup>7</sup>Differential privacy was popularised by [19]. A useful lay introduction is provided in [37]

<sup>8</sup>Pure differential privacy would mean that the analysis is *exactly* the same; in practice there is a small margin of flexibility that is allowed.

vulnerable to model inversion. In some cases, such attacks have been shown to succeed in part even in the absence of overfitting [52].

#### **4.5 Storage limitation principle**

Relatedly, another data protection principle, that of ‘storage limitation’, applies in this case. Storage limitation means that data should be kept for no longer than it is necessary for the purposes for which it is processed. As training data may need to be discarded as time goes on to meet this obligation, so might models. Some techniques from the field of concept drift adaptation could be useful here. This is the domain of research which looks at how to understand and model changing phenomena in machine learning systems. For example, [25] describe a variety of methods used to limit the use of older data in machine learning systems by both managing data used in machine learning and gradually forgetting data within a model. These methods are primarily used to better model phenomena today, particularly where old correlations may now be irrelevant, however similar efforts are likely to be required for any systems for which model inversion is not a readily remediable vulnerability.

### **5 Discussion and conclusion**

In this paper, we have outlined how recent confidentiality attacks upon machine learning systems, including model inversion and membership inference, interplay with data protection law. Where models are vulnerable to such attacks, they gain an additional dimension—not only an analytic product potentially protected by intellectual property rights, but also a set of personal data, conceptually close to the idea of ‘pseudonymisation’ in the GDPR. We illustrated a selection of consequences that could flow from this new classification that are relevant to discussions of the governance of machine learned models which are being traded and transferred between data controllers. These include those of direct utility to data subjects, such as information, erasure and objection rights, and overarching obligations relevant to data controllers, such as security and storage limitation provisions. Seeing models as personal data could serve to re-balance or at least disrupt the power relations between those holding models and those whose data is used to train them.

There is however reason for caution amidst the promise. Whilst potentially enabling useful provisions, requiring what is essentially a security vulnerability in order to trigger rights and obligations is disconnected and arbitrary. Models not amenable to model inversion might still be models individuals wish to know the origin or destination of or wish to have themselves or their group erased from. We suspect many situations of problematic profiling will not be vulnerable to model inversion, and therefore not governable with this approach. Furthermore,



where some model inversion is possible, technical challenges will make it difficult for data subjects and regulators to prove the leakiness. In many cases, while a companion dataset that may enable such an attack might exist, potentially on shady markets, it will not be held by the data subject, auditor nor the regulator, and thus the risk will be difficult to assess even if full model access is provided. The setup appears to further burden every stakeholder apart from the model, and it remains questionable whether downstream governance provisions in general, such as rights after model training, are the best way to deal with algorithmic harms at all [20].

Data protection provisions were not designed with the profiling and inference capabilities we see today. It is becoming clearer that many socio-technical challenges presented by machine learning and algorithmic systems more broadly are not wholly dealt with using the provisions in regulations such as the GDPR, which are the result of a slow evolution in definitions and concerns. On one hand, this has created a set of principles that are desirable and robust for many purposes. On the other hand, it is showing that seeing data protection as an omnibus governing regime for all data-driven issues is misguided. In this paper, we illustrate this by considering model inversion as a regulatory experiment—one that is potentially realistic, although far from advisable as a foundation for future algorithmic governance. We argue it provides a hook to probe whether or not it is appropriate these rights and obligations extend to analytic products themselves, and what the consequences of this development might be. Some consequences, such as the mapping of the provenance of trained models, seem potentially useful in oversight of increasingly influential systems. Others, such as the right to object and restrict processing, seem at tension with the very notion of model inversion. As the domain of personal data expands, it is important to recognise that while its scope is wide, its scope of effective and enforceable governance might not be. Pushing the boundaries of data protection law in light of new technologies serves as a useful and clarifying force in the continuous task of the coming decades in better applying existing law, and developing regulatory regimes around societal issues they fail to deal with in practice.

## References

- [1] Martin Abadi et al. “On the Protection of Private Information in Machine Learning Systems: Two Recent Approaches”. In: *Proceedings of the 30th IEEE Computer Security Foundations Symposium, August 21-25, 2017, Santa Barbara, CA, USA*. 2017.
- [2] Article 29 Working Party. *Guidelines on transparency under Regulation 2016/679 (WP260)*. 2018.

- [3] Giuseppe Ateniese et al. “Hacking smart machines with smarter ones: How to extract meaningful data from machine learning classifiers”. In: *IJSN* 10.3 (2015), p. 137.
- [4] Jane Bambauer, Krishnamurty Muralidhar, and Rathindra Sarathy. “Fool’s Gold: An Illustrated Critique of Differential Privacy”. In: *Vand. J. Ent. & Tech. L.* 16 (2013), p. 701.
- [5] David Banisar. *The right to information and privacy: balancing rights and managing conflicts*. Washington, DC: World Bank, 2011.
- [6] Debjanee Barua. “A time to remember, a time to forget: User controlled, Scalable, Life long user modelling”. PhD thesis. The University of Sydney, 2016.
- [7] Helen L Bear, Stephen J Cox, and Richard W Harvey. “Speaker-independent machine lip-reading with speaker-dependent viseme classifiers”. In: *The 1st Joint Conference on Facial Analysis, Animation, and Auditory-Visual Speech Processing (FAAVSP-2015), Vienna, Austria, September 11-13, 2015*. 2015, pp. 190–195.
- [8] Reuben Binns et al. “Third Party Tracking in the Mobile Ecosystem”. In: *Proceedings of the 10th ACM Conference on Web Science*. 2018.
- [9] Tim Bradshaw. *Apple plays catch-up with iMessage emojis*. <https://www.ft.com/content/1d9a1aa6-31dc-11e6-bda0-04585c31b153>. Accessed: 2017-10-11. June 2016.
- [10] Joseph A Calandrino et al. ““You Might Also Like:” Privacy Risks of Collaborative Filtering”. In: *Security and Privacy (SP), 2011 IEEE Symposium on*. IEEE. 2011, pp. 231–246.
- [11] Y Cao and J Yang. “Towards Making Systems Forget with Machine Unlearning”. In: *2015 IEEE Symposium on Security and Privacy*. 2015, pp. 463–480.
- [12] Nicholas Carlini et al. “The Secret Sharer: Measuring Unintended Neural Network Memorization & Extracting Secrets”. In: *arXiv preprint arXiv:1802.08232* (2018).
- [13] Kamalika Chaudhuri, Claire Monteleoni, and Anand D Sarwate. “Differentially private empirical risk minimization”. In: *J. Mach. Learn. Res.* 12 (2011), pp. 1069–1109.
- [14] Andy Crabtree et al. “Enabling the new economic actor: data protection, the digital economy, and the Databox”. In: *Personal and Ubiquitous Computing* 20.6 (2016), pp. 947–957.
- [15] Jon Crowcroft et al. “Unclouded Vision”. In: *Distributed Computing and Networking*. Ed. by Marcos K. Aguilera et al. Berlin: Springer, 2011, pp. 29–40. ISBN: 978-3-642-17679-1.

- [16] Sari Depreeuw and Katja de Vries. *Deliverable 3.11: Profile transparency, trade secrets and intellectual property rights in OSNs*. User Empowerment for Enhanced Online Presence Management (USEMP) project, 2016.
- [17] Cory Doctorow. “Data protection in the EU: the certainty of uncertainty”. In: *The Guardian* (June 2013).
- [18] C Dwork. “Differential privacy: A survey of results”. In: *International Conference on Theory and Applications of Models of Computation*. Dordrecht: Springer, 2008.
- [19] Cynthia Dwork et al. “Calibrating Noise to Sensitivity in Private Data Analysis”. In: *Theory of Cryptography: Third Theory of Cryptography Conference, TCC 2006, New York, NY, USA, March 4-7, 2006. Proceedings*. Ed. by Shai Halevi and Tal Rabin. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 265–284.
- [20] Lilian Edwards and Michael Veale. “Slave to the Algorithm? Why a ‘Right to an Explanation’ is Probably Not the Remedy You are Looking For”. In: *Duke Law Technol. Rev.* 16.1 (2017), pp. 18–84.
- [21] Farhan et al. “Behavior vs. introspection: refining prediction of clinical depression via smartphone sensing data”. In: *2016 IEEE Wireless Health (WH)*. Oct. 2016, pp. 1–8.
- [22] Matt Fredrikson, Somesh Jha, and Thomas Ristenpart. “Model Inversion Attacks that Exploit Confidence Information and Basic Countermeasures”. In: *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*. ACM, 2015, pp. 1322–1333.
- [23] Matthew Fredrikson et al. “Privacy in Pharmacogenetics: An End-to-End Case Study of Personalized Warfarin Dosing”. In: *USENIX Security Symposium*. 2014, pp. 17–32.
- [24] Antonella Galetta and Paul de Hert. “Exercising Access Rights in Belgium”. In: *The Unaccountable State of Surveillance: Exercising Access Rights in Europe*. Ed. by Clive Norris et al. Cham: Springer, 2017, pp. 77–108.
- [25] J Gama et al. “A survey on concept drift adaptation”. In: *ACM Computing Surveys* 1.1 (2013).
- [26] Google. *Transparency Report*. Accessed: 2017-10-10.
- [27] Mireille Hildebrandt. “Profiling and the rule of law”. en. In: *Identity in the Information Society* 1.1 (Dec. 2008), pp. 55–70.
- [28] Mireille Hildebrandt. *Smart technologies and the end(s) of law*. Cheltenham, UK: Edward Elgar, 2015.
- [29] Zhanglong Ji, Zachary C Lipton, and Charles Elkan. “Differential Privacy and Machine Learning: a Survey and Review”. In: *arXiv preprint* (2014).
- [30] S Kasiviswanathan et al. “What Can We Learn Privately?” In: *SIAM J. Comput.* 40.3 (Jan. 2011), pp. 793–826.

- [31] René Mahieu, Hadi Asghari, and Michel van Eeten. “Collectively Exercising the Right of Access: Individual Effort, Societal Effect”. In: *GigaNet (Global Internet Governance Academic Network) Annual Symposium, Geneva, Switzerland*. 2017. URL: <https://ssrn.com/abstract=3107292>.
- [32] Toni Mastelic et al. “Cloud Computing: Survey on Energy Efficiency”. In: *ACM CSUR* 47.2 (2015), p. 33.
- [33] Frank McSherry. “Statistical inference considered harmful”. In: (June 2016). URL: <https://github.com/frankmcsherry/blog/blob/master/posts/2016-06-14.md>.
- [34] Ilya Mironov. “On Significance of the Least Significant Bits for Differential Privacy”. In: *Proceedings of the 2012 ACM Conference on Computer and Communications Security*. CCS ’12. New York, NY, USA: ACM, 2012, pp. 650–661.
- [35] Brent Mittelstadt. “From Individual to Group Privacy in Big Data Analytics”. In: *Philos. Technol.* (Feb. 2017), pp. 1–20.
- [36] Krish Muralidhar and Rathindra Sarathy. “Does Differential Privacy Protect Terry Gross’ Privacy?” In: *Privacy in Statistical Databases: UNESCO Chair in Data Privacy, International Conference, PSD 2010, Corfu, Greece, September 22–24, 2010. Proceedings*. Ed. by Josep Domingo-Ferrer and Emmanouil Magkos. Berlin, Heidelberg: Springer Berlin Heidelberg, 2010, pp. 200–209.
- [37] Kobbi Nissim et al. *Differential Privacy: A Primer for a Non-technical Audience*. Cambridge, MA: A product of the “Bridging Privacy Definitions” working group, part of the Privacy Tools for Sharing Research Data project at Harvard University., 2017.
- [38] Paul Ohm. “Broken promises of privacy: Responding to the surprising failure of anonymization”. In: *UCLA L. Rev.* 57 (2009), p. 1701.
- [39] Andrew Orlowski. *Apple pollutes data about you to protect your privacy. But it might not be enough*. [https://www.theregister.co.uk/2016/06/21/apple\\_pollutes\\_your\\_data\\_differential\\_privacy/](https://www.theregister.co.uk/2016/06/21/apple_pollutes_your_data_differential_privacy/). Accessed: 2017-10-11. June 2016.
- [40] Nadezdha Purtova. “The law of everything. Broad concept of personal data and future of EU data protection law”. In: *Law, Innovation, and Technology* 10 (1 2018).
- [41] Apostolos Pyrgelis, Carmela Troncoso, and Emiliano De Cristofaro. “Knock Knock, Who’s There? Membership Inference on Aggregate Location Data”. In: *arXiv preprint* (2017).
- [42] Andrew D Selbst and Julia Powles. “Meaningful information and the right to explanation”. In: *International Data Privacy Law* 7.4 (2017), pp. 233–242.
- [43] Reza Shokri et al. “Membership Inference Attacks Against Machine Learning Models”. In: *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, May 2017, pp. 3–18.

- [44] Jatinder Singh, Jennifer Cobbe, and Chris Norval. “Decision Provenance: Capturing data flow for accountable systems”. In: *arXiv preprint arXiv:1804.05741* (2018).
- [45] Congzheng Song, Thomas Ristenpart, and Vitaly Shmatikov. “Machine Learning Models that Remember Too Much”. In: *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS ’17)*. New York, New York, USA: ACM Press, 2017, pp. 587–601.
- [46] Linnet Taylor, Luciano Floridi, and Bart van der Sloot. *Group Privacy*. en. Springer, Jan. 2017.
- [47] Florian Tramèr et al. “Stealing Machine Learning Models via Prediction APIs”. In: *USENIX Security Symposium*. 2016, pp. 601–618.
- [48] Max Van Kleek et al. “Better the Devil You Know: Exposing the Data Sharing Practices of Smartphone Apps”. In: *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems (CHI’17)*. 2017, pp. 5208–5220.
- [49] Michael Veale, Reuben Binns, and Jef Ausloos. “When data protection by design and data subject rights clash”. In: *International Data Privacy Law* (2018).
- [50] Anton Vedder. “KDD: The challenge to individualism”. en. In: *Ethics Inf. Technol.* 1.4 (Dec. 1999), pp. 275–281.
- [51] Sandra Wachter, Brent Mittelstadt, and Luciano Floridi. “Why a right to explanation of automated decision-making does not exist in the General Data Protection Regulation”. In: *International Data Privacy Law* 7.2 (2017), pp. 76–99.
- [52] Samuel Yeom, Matt Fredrikson, and Somesh Jha. “The Unintended Consequences of Overfitting: Training Data Inference Attacks”. In: *arXiv preprint arXiv:1709.01604* (2017).
- [53] Zhi-Hua Zhou. “Learnware: on the future of machine learning”. In: *Front. Comput. Sci* 10.4 (2016), pp. 589–590.